

# Traffic\_Sign\_Classifier

March 26, 2017

```
In [135]: # Load pickled data
import pickle, os

# TODO: Fill this in based on where you saved the training and testing data

training_file = os.path.join(os.getcwd(), './train.p')
validation_file = os.path.join(os.getcwd(), './valid.p')
testing_file = os.path.join(os.getcwd(), './test.p')

with open(training_file, mode='rb') as f:
    train = pickle.load(f)
with open(validation_file, mode='rb') as f:
    valid = pickle.load(f)
with open(testing_file, mode='rb') as f:
    test = pickle.load(f)

X_train, y_train = train['features'], train['labels']
X_valid, y_valid = valid['features'], valid['labels']
X_test, y_test = test['features'], test['labels']

In [136]: import matplotlib.pyplot as plt
import matplotlib.image as mpimg
import cv2
%matplotlib inline

import numpy as np
from sklearn.utils import shuffle
# Number of training examples
n_train = X_train.shape[0]
n_valid = X_valid.shape[0]
# Number of testing examples.
n_test = X_test.shape[0]

# What's the shape of an traffic sign image?
image_shape = X_train.shape[1:3]

# TODO: How many unique classes/labels there are in the dataset.
```

```

n_classes = np.unique(y_train).size

print("Number of training examples =", n_train)
print("Number of testing examples =", n_test)
print("Number of validation examples = {0}".format(n_valid))
print("Image data shape =", image_shape)
print("Number of classes =", n_classes)

```

```

Number of training examples = 34799
Number of testing examples = 12630
Number of validation examples = 4410
Image data shape = (32, 32)
Number of classes = 43

```

```

In [137]: # Signnames CSV loaded to make sure sign names stay attached with data du
import pandas as pd
sign_names = pd.read_csv('signnames.csv', index_col=0)
sign_names['SignName'][0]

```

```

Out[137]: 'Speed limit (20km/h)'

```

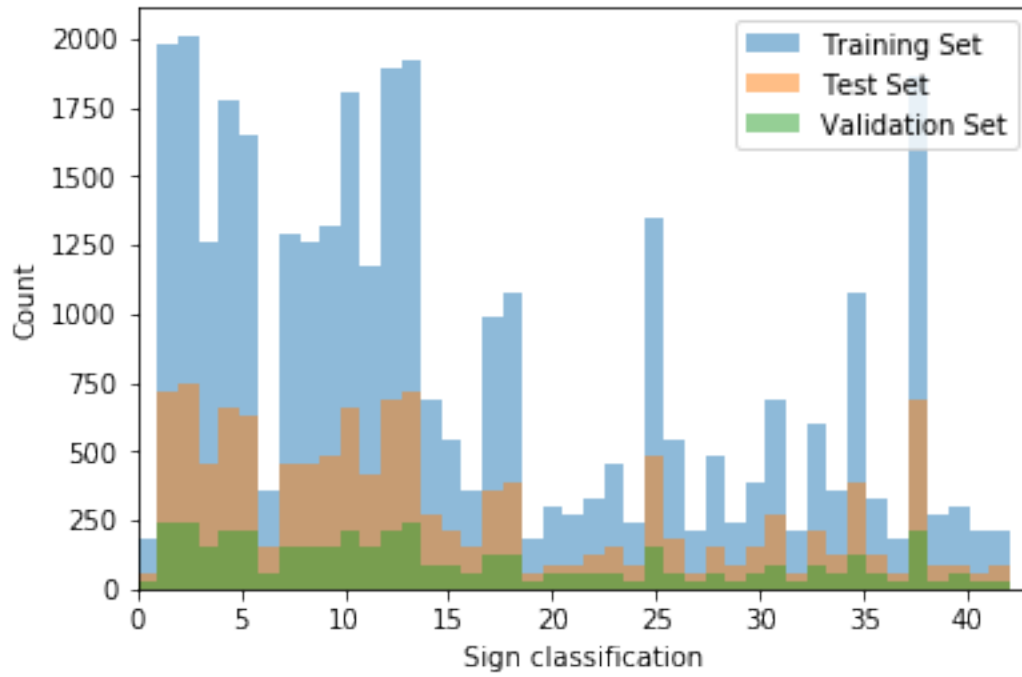
## 1 Data Exploration

- Data plotted to show relative number of test examples for each sign in each of training, validation, and test set.
- Datasets are not represented equally.
- Images and names of signs shown from testing and validation set

```

In [138]: plt.hist(y_train, n_classes, alpha=0.5, label='Training Set')
plt.hist(y_test, n_classes, alpha=0.5, label='Test Set')
plt.hist(y_valid, n_classes, alpha=0.5, label='Validation Set')
plt.xlabel('Sign classification')
plt.ylabel('Count')
plt.xlim(xmin=0, xmax=n_classes)
plt.legend(loc='upper right')
plt.show()

```



```
In [148]: X_train, y_train = shuffle(X_train, y_train)
fig = plt.figure(figsize=(15,18))
fig.suptitle('Exploring data')
for i in range(1,26):
    ax = plt.subplot(5,5,i)
    ax.set_title(sign_names['SignName'][y_train[i]])
    plt.imshow(X_train[i])
plt.show()
```

## Exploring data



```
In [147]: X_valid, y_valid = shuffle(X_valid, y_valid)
fig = plt.figure(figsize=(15,18))
fig.suptitle('Exploring data')
for i in range(1,26):
    ax = plt.subplot(5,5,i)
```

```

ax.set_title(sign_names['SignName'][y_valid[i]])
plt.imshow(X_valid[i])
plt.show()

```

Exploring data



## 2 Data Amplification

1. Rotation
2. Histogram equalization (enhance contrast)
3. Contrast limited adaptive histogram equalization image contrast
4. Conorm (enhance edge with difference of gaussian)

### 2.1 Rotation +/- 5 degrees

- Traffic signs may not be oriented the same as in the training images.
- Images show greater rotation for visual clarity, smaller angles expected in real world
- Border mode changed from constant (black) to replicate. I expect the large shift in color from constant border could impact gradients at the edges and confuse the model.

```
In [149]: def rotate_images(X, deg):
    rows, cols = X.shape[1], X.shape[2]
    M_pos = cv2.getRotationMatrix2D((rows//2, cols//2), deg, 1)
    M_neg = cv2.getRotationMatrix2D((rows//2, cols//2), -1*deg, 1)
    X_rot_pos = np.copy(X)
    X_rot_neg = np.copy(X)
    for i in range(X.shape[0]):
        X_rot_pos[i] = cv2.warpAffine(X[i], M_pos, (rows, cols), borderMode=cv2.BORDER_REPLICATE)
        X_rot_neg[i] = cv2.warpAffine(X[i], M_neg, (rows, cols), borderMode=cv2.BORDER_REPLICATE)
    print('Dataset rotated +/-{0} degrees.'.format(deg))
    return X_rot_pos, X_rot_neg
```

```
In [150]: def rotate_images_black(X, deg):
    rows, cols = X.shape[1], X.shape[2]
    M_pos = cv2.getRotationMatrix2D((rows//2, cols//2), deg, 1)
    M_neg = cv2.getRotationMatrix2D((rows//2, cols//2), -1*deg, 1)
    X_rot_pos = np.copy(X)
    X_rot_neg = np.copy(X)
    for i in range(X.shape[0]):
        X_rot_pos[i] = cv2.warpAffine(X[i], M_pos, (rows, cols), borderMode=cv2.BORDER_CONSTANT, borderValue=(0,0,0))
        X_rot_neg[i] = cv2.warpAffine(X[i], M_neg, (rows, cols), borderMode=cv2.BORDER_CONSTANT, borderValue=(0,0,0))
    print('Dataset rotated +/-{0} degrees.'.format(deg))
    return X_rot_pos, X_rot_neg
```

#### 2.1.1 Replicated border

```
In [151]: degrees = 15
    X_pos, X_neg = rotate_images(X_train, degrees)
    X_train, y_train, X_pos, X_neg = shuffle(X_train, y_train, X_pos, X_neg)
    fig = plt.figure(figsize=(12,8))
    fig.suptitle('+/-{0} degree rotation'.format(degrees))

    ax1 = plt.subplot(231)
    ax1.set_title(sign_names['SignName'][y_train[0]])
    plt.imshow(X_pos[0])
```

```

ax2 = plt.subplot(232)
ax2.set_title(sign_names['SignName'][y_train[0]])
plt.imshow(X_train[0])

ax3 = plt.subplot(233)
ax3.set_title(sign_names['SignName'][y_train[0]])
plt.imshow(X_neg[0])


ax4 = plt.subplot(234)
ax4.set_title(sign_names['SignName'][y_train[1]])
plt.imshow(X_pos[1])

ax5 = plt.subplot(235)
ax5.set_title(sign_names['SignName'][y_train[1]])
plt.imshow(X_train[1])

ax6 = plt.subplot(236)
ax6.set_title(sign_names['SignName'][y_train[1]])
plt.imshow(X_neg[1])
plt.show()

```

Dataset rotated  $\pm 15$  degrees.



## 2.1.2 Constant border

```
In [152]: degrees = 15
          X_pos, X_neg = rotate_images_black(X_train, degrees)
          fig = plt.figure(figsize=(12,8))
          fig.suptitle(' +/-{0} degree rotation'.format(degrees))

          ax1 = plt.subplot(231)
          ax1.set_title(sign_names['SignName'][y_train[0]])
          plt.imshow(X_pos[0])

          ax2 = plt.subplot(232)
          ax2.set_title(sign_names['SignName'][y_train[0]])
          plt.imshow(X_train[0])

          ax3 = plt.subplot(233)
          ax3.set_title(sign_names['SignName'][y_train[0]])
          plt.imshow(X_neg[0])

          ax4 = plt.subplot(234)
          ax4.set_title(sign_names['SignName'][y_train[1]])
          plt.imshow(X_pos[1])


          ax5 = plt.subplot(235)
          ax5.set_title(sign_names['SignName'][y_train[1]])
          plt.imshow(X_train[1])

          ax6 = plt.subplot(236)
          ax6.set_title(sign_names['SignName'][y_train[1]])
          plt.imshow(X_neg[1])
          plt.show()
```

Dataset rotated +/-15 degrees.



+/-15 degree rotation



## 2.2 Increase image contrast

- Histogram equalization is used to increase contrast in image
- Color histogram equalization achieved by equalizing Y channel of YCrCb colorspace

```
In [153]: def amplify_contrast(X):
           X_eq = np.copy(X)
           for i in range(X.shape[0]):
               YCrCb = cv2.cvtColor(X[i], cv2.COLOR_RGB2YCrCb)
               Y, Cr, Cb = cv2.split(YCrCb)
               Y_eq = cv2.equalizeHist(Y)
               YCrCb_eq = cv2.merge((Y_eq, Cr, Cb))
               X_eq[i] = cv2.cvtColor(YCrCb_eq, cv2.COLOR_YCrCb2RGB)
           print('Dataset histogram equalized')
           return X_eq

In [156]: X_histeq = amplify_contrast(X_train)
           X_train, y_train, X_histeq = shuffle(X_train, y_train, X_histeq)
           fig = plt.figure(figsize=(8,8))
           fig.suptitle('Increase contrast histogram equalization')
           ax1 = plt.subplot(221)
```

```
ax1.set_title(sign_names['SignName'][y_train[0]])
plt.imshow(X_train[0])


ax2 = plt.subplot(222)
ax2.set_title(sign_names['SignName'][y_train[0]])
plt.imshow(X_histeq[0])

ax3 = plt.subplot(223)
ax3.set_title(sign_names['SignName'][y_train[1]])
plt.imshow(X_train[1])

ax4 = plt.subplot(224)
ax4.set_title(sign_names['SignName'][y_train[1]])
plt.imshow(X_histeq[1])
plt.show()
```

Dataset histogram equalized

## Increase contrast histogram equalization



### 2.3 Increase image contrast (Area)

- Adaptive contrast limited histogram equalization is used to increase contrast in image
- Different from histogram equalization in that it looks at blocks of pixels.
- Color histogram equalization achieved by equalizing Y channel of YCrCb colorspace

```
In [157]: def clahe_contrast(X):  
           X_clahe = np.copy(X)  
           clahe = cv2.createCLAHE(tileGridSize=(4,4))  
           for i in range(X.shape[0]):
```

```

        YCrCb = cv2.cvtColor(X[i], cv2.COLOR_RGB2YCrCb)
        Y, Cr, Cb = cv2.split(YCrCb)
        Y_clahe = clahe.apply(Y)
        YCrCb_clahe = cv2.merge((Y_clahe, Cr, Cb))
        X_clahe[i] = cv2.cvtColor(YCrCb_clahe, cv2.COLOR_YCrCb2RGB)
    print('Dataset histogram equalized clahe')
    return X_clahe

```

```

In [165]: X_clahe = clahe_contrast(X_train)
          X_train, y_train, X_clahe = shuffle(X_train, y_train, X_clahe)
          fig = plt.figure(figsize=(8,8))
          fig.suptitle('Increase contrast CLAHE')

          ax1 = plt.subplot(221)
          ax1.set_title(sign_names['SignName'][y_train[0]])
          plt.imshow(X_train[0])

          ax2 = plt.subplot(222)
          ax2.set_title(sign_names['SignName'][y_train[0]])
          plt.imshow(X_clahe[0])


          ax3 = plt.subplot(223)
          ax3.set_title(sign_names['SignName'][y_train[1]])
          plt.imshow(X_train[1])

          ax4 = plt.subplot(224)
          ax4.set_title(sign_names['SignName'][y_train[1]])
          plt.imshow(X_clahe[1])
          plt.show()

```

Dataset histogram equalized clahe

## Increase contrast CLAHE



### 2.4 Contrast Normalization

- Use a difference of gaussian to amplify edges
- Obtained by convolving a difference of gaussians on the image and adding that to the image
- Based on [Sermanet Lecun 2011](#)
- Can lead to artifacts in color images. I don't understand why that is, so I'm not including this algorithm

```
In [166]: def contrast_normalization(X):  
          X_conorm = np.copy(X)
```

```

gauss_1 = cv2.getGaussianKernel(5, 1) + np.transpose(cv2.getGaussianK
gauss_2 = cv2.getGaussianKernel(5, 3) + np.transpose(cv2.getGaussianK
diff_gauss = gauss_1 - gauss_2
for i in range(X.shape[0]):
    X_conorm[i] = X[i] + cv2.filter2D(X[i], -1, diff_gauss)
print('Contrast normalized')
return X_conorm

```

```

In [183]: X_conorm = contrast_normalization(X_train)
X_train, y_train, X_conorm = shuffle(X_train, y_train, X_conorm)
fig = plt.figure(figsize=(8,8))
fig.suptitle('Contrast normalization')

ax1 = plt.subplot(221)
ax1.set_title(sign_names['SignName'][y_train[0]])
plt.imshow(X_train[0])

ax2 = plt.subplot(222)
ax2.set_title(sign_names['SignName'][y_train[0]])
plt.imshow(X_conorm[0])


ax3 = plt.subplot(223)
ax3.set_title(sign_names['SignName'][y_train[1]])
plt.imshow(X_train[1])

ax4 = plt.subplot(224)
ax4.set_title(sign_names['SignName'][y_train[1]])
plt.imshow(X_conorm[1])
plt.show()

```

Contrast normalized

## Contrast normalization



```
In [184]: # Note contrast normalization removed. It doesn't look like it's working
def amplify_dataset(X, y):
    X_rot_pos, X_rot_neg = rotate_images(X, 5)
    X = np.concatenate((X, X_rot_neg, X_rot_pos), axis=0)
    y = np.concatenate((y, y, y), axis=0)

    X_histeq = amplify_contrast(X)
    X_clahe = clahe_contrast(X)
    X_conorm = contrast_normalization(X)
    X = np.concatenate((X, X_histeq, X_clahe), axis=0)
```

```

y = np.concatenate((y, y, y), axis=0)
print('Number of examples increased to {}'.format(X.shape[0]))
return X, y

```

```
In [185]: X_train, y_train = amplify_dataset(X_train, y_train)
```

```

Dataset rotated +/-5 degrees.
Dataset histogram equalized
Dataset histogram equalized clahe
Contrast normalized
Number of examples increased to 313191

```

### 3 Preprocess Data

- Performed on training, validation and test data.
- Grayscale images to reduce the number of parameters.
- [Sermanet Lecun 2011](#) found color didn't improve accuracy a lot. Prior models run with color agreed with this.
- Data is normalized by subtracting the mean image and normalizing following [common best practices](#) in order to keep our features in a consistent range. This will reduce the likelihood of our gradients getting out of control through vanishing gradient / saturating neurons in the network

```

In [186]: def grayscale(X):
           X_gray = np.copy(X[:, :, :, 0])
           for i in range(X.shape[0]):
               X_gray[i] = cv2.cvtColor(X[i], cv2.COLOR_RGB2GRAY)
           return X_gray

```

```

In [187]: #Normalize the data
           def normalize(X):
               X = X.astype('float64') - np.mean(X, axis=0)
               X /= np.std(X, axis=0)
               return X

```

```

In [188]: def preprocess(X):
           X_gray = grayscale(X)
           X_norm = normalize(X_gray)
           return X_norm

```

```

In [189]: X_train = preprocess(X_train)
           X_valid = preprocess(X_valid)
           X_test = preprocess(X_test)

```

#### 3.1 Model Architecture

1. Follows similar structure to [Ciresan et al 2012](#)



2. Consists of 3 sets of two convolutional layers with ReLU activation functions and a pooling layer
3. Same padding used to limit reduction in image size since starting image size was smaller than paper
4. A fully connected layer is followed by a dropout layer.
5. Dropout is used for regularization to reduce overfitting of the model.
6. Two more hidden layers are used before the final output layer
7. The exact shape of every layer is in the picture below

Input	32x32x1
Conv-32	32x32x32
Conv-32	32x32x32
Pool	16x16x32
Conv-64	16x16x64
Conv-64	16x16x64
Pool	8x8x64
Conv-128	8x8x128
Conv-128	8x8x128
Pool	4x4x128
FC	2048
FC	400
FC	100
FC	43
Softmax	

```
In [13]: import tensorflow as tf
         print(tf.__version__)
```

1.0.0

```

In [17]: def weight_variable(shape):
          return tf.Variable(tf.truncated_normal(shape, mean=0.01, stddev=0.1))
          def bias_variable(shape):
              return tf.Variable(tf.zeros(shape=shape))

In [18]: def conv2d(x, W):
          return tf.nn.conv2d(x, W, strides=[1,1,1,1], padding='VALID')
          def conv2d_same(x, W):
              return tf.nn.conv2d(x, W, strides=[1,1,1,1], padding='SAME')
          def max_pool(x):
              return tf.nn.max_pool(x, ksize=[1,2,2,1],strides=[1,2,2,1],
                                     padding='VALID')

```

### 3.2 Train, Validate and Test the Model

A validation set can be used to assess how well the model is performing. A low accuracy on the training and validation sets imply underfitting. A high accuracy on the training set but low accuracy on the validation set implies overfitting.

1. Learning rate set to 1e-4. The learning rate of Adam decays over time and adjusts based on the gradient of the variables with momentum.
2. Adam Optimizer adjusts the learning rate using the [Adam algorithm](#) This optimizer is derived from Adagrad an adaptive learning algorithm. The algorithm monotonically reduces its learning rate over time. RMSProp improved on this by using a moving average of gradients to reduce the aggressiveness of Adagrad. Adam improved on RMSProp by adding momentum. The default momentum and decay rate are used from tensorflow.
3. A batch size of 64 is a function of using a large VGG-like architecture and having an older GPU. The first LeNet-like architecture had a batch size of 512.
4. Dropout parameter of 0.5 is used as a default because it works [Srivastava et al 2014](#) supposedly because it maximizes the regularization value of the dropout layer.
5. Based on VGGNet
6. The number of epochs was set to 1000 originally. This is because it was a large number at which point it seemed like the loss had reached a plateau. When training multiple networks 500 was used because it took a long time and graphs of validation accuracy of the prior run had plateaued around then.

```

In [19]: def SketchNet(input_layer, dropout, num_classes=43):
          ''' Neural Network modeled on VGG Net and scaled down for my machine.
          Params: input_layer:  4D tensor holding 32x32 grayscale images
                  dropout:      Fraction of neurons to kill during training using
                  num_classes:  Number of discrete classes this network will pre
          '''
          # depth = [1, 6, 16, 400, 160, 96, num_classes]
          input_layer = tf.reshape(x, [-1,32,32, 1])
          print('Layer 0 Init: Shape {}'.format(input_layer.get_shape()))

          F_w_1 = weight_variable([5,5,1,32])

```

```

F_b_1 = bias_variable([32])
conv_1 = tf.nn.relu(conv2d_same(input_layer, F_w_1) + F_b_1)
print('Layer 1 Convolution: Shape {}'.format(conv_1.get_shape()))
F_w_2 = weight_variable([3,3,32,32])
F_b_2 = bias_variable([32])
conv_2 = tf.nn.relu(conv2d_same(conv_1, F_w_2) + F_b_2)
print('Layer 2 Convolution: Shape {}'.format(conv_2.get_shape()))
pool_1 = max_pool(conv_2)
print('Layer 3 Pooling: Shape {}'.format(pool_1.get_shape()))

F_w_3 = weight_variable([3,3,32,64])
F_b_3 = bias_variable([64])
conv_3 = tf.nn.relu(conv2d_same(pool_1, F_w_3) + F_b_3)
print('Layer 4 Convolution: Shape {}'.format(conv_3.get_shape()))
F_w_4 = weight_variable([3,3,64,64])
F_b_4 = bias_variable([64])
conv_4 = tf.nn.relu(conv2d_same(conv_3, F_w_4) + F_b_4)
print('Layer 5 Convolution: Shape {}'.format(conv_4.get_shape()))
pool_2 = max_pool(conv_4)
print('Layer 6 Pooling: Shape {}'.format(pool_2.get_shape()))

F_w_5 = weight_variable([3,3,64,128])
F_b_5 = bias_variable([128])
conv_5 = tf.nn.relu(conv2d_same(pool_2, F_w_5) + F_b_5)
print('Layer 7 Convolution: Shape {}'.format(conv_5.get_shape()))
F_w_6 = weight_variable([3,3,128,128])
F_b_6 = bias_variable([128])
conv_6 = tf.nn.relu(conv2d_same(conv_5, F_w_6) + F_b_6)
print('Layer 8 Convolution: Shape {}'.format(conv_6.get_shape()))
pool_3 = max_pool(conv_6)
print('Layer 9 Pooling: Shape {}'.format(pool_3.get_shape()))

dense_1 = tf.contrib.layers.flatten(pool_3)
print('Layer 10 Flatten: Shape {}'.format(dense_1.get_shape()))

drop_1 = tf.nn.dropout(dense_1, dropout)

F_w_7 = weight_variable([2048, 400])
F_b_7 = bias_variable([400])
dense_2 = tf.nn.relu(tf.matmul(drop_1, F_w_7) + F_b_7)
print('Layer 11 Matrix Multiplication: Shape {}'.format(dense_2.get_shape()))

F_w_8 = weight_variable([400, 100])
F_b_8 = bias_variable([100])
dense_3 = tf.nn.relu(tf.matmul(dense_2, F_w_8) + F_b_8)
print('Layer 12 Matrix Multiplication: Shape {}'.format(dense_3.get_shape()))

F_w_9 = weight_variable([100, n_classes])

```

```

        F_b_9 = bias_variable([n_classes])
        logits = tf.nn.relu(tf.matmul(dense_3, F_w_9) + F_b_9)
        print('Final Layer Matrix Multiplication: Shape {}'.format(logits.get_shape().as_list()))

    return logits

In [20]: tf.reset_default_graph()
x = tf.placeholder(tf.float32, (None, 32, 32))
y = tf.placeholder(tf.int32, (None))
keep_prob = tf.placeholder(tf.float32)
# x = tf.constant(X_train_rot_gray)
one_hot_y = tf.one_hot(y, n_classes)

rate = 0.0001
pred = SketchNet(x, keep_prob)
cross_entropy = tf.nn.softmax_cross_entropy_with_logits(logits=pred,
                                                         labels=one_hot_y)

loss_operation = tf.reduce_mean(cross_entropy)
optimizer = tf.train.AdamOptimizer(learning_rate=rate)
training_operation = optimizer.minimize(loss_operation)
correct_pred = tf.equal(tf.argmax(pred, 1), tf.argmax(one_hot_y, 1))
accuracy_op = tf.reduce_mean(tf.cast(correct_pred, tf.float32))

Layer 0 Init: Shape (?, 32, 32, 1)
Layer 1 Convolution: Shape (?, 32, 32, 32)
Layer 2 Convolution: Shape (?, 32, 32, 32)
Layer 3 Pooling: Shape (?, 16, 16, 32)
Layer 4 Convolution: Shape (?, 16, 16, 64)
Layer 5 Convolution: Shape (?, 16, 16, 64)
Layer 6 Pooling: Shape (?, 8, 8, 64)
Layer 7 Convolution: Shape (?, 8, 8, 128)
Layer 8 Convolution: Shape (?, 8, 8, 128)
Layer 9 Pooling: Shape (?, 4, 4, 128)
Layer 10 Flatten: Shape (?, 2048)
Layer 11 Matrix Multiplication: Shape (?, 400)
Layer 12 Matrix Multiplication: Shape (?, 100)
Final Layer Matrix Multiplication: Shape (?, 43)

In [89]: def evaluate(X_data, y_data, sess):
    num_examples = len(X_data)
    total_accuracy = 0
    for offset in range(0, num_examples, BATCH_SIZE):
        batch_x = X_data[offset:offset+BATCH_SIZE]
        batch_y = y_data[offset:offset+BATCH_SIZE]
        accuracy = sess.run(accuracy_op, feed_dict={x: batch_x,
                                                    y: batch_y,

```

```

                                                    keep_prob: 1.0})
    total_accuracy += (accuracy * len(batch_x))
    return total_accuracy / num_examples

```

### 3.2.1 Training multiple networks

The below training sought to train multiple networks for testing an ensemble network. The same code works to train a single network.

```

In [30]: for i in range(7,10):
    tf.reset_default_graph()
    x = tf.placeholder(tf.float32, (None, 32, 32))
    y = tf.placeholder(tf.int32, (None))
    keep_prob = tf.placeholder(tf.float32)
    # x = tf.constant(X_train_rot_gray)
    one_hot_y = tf.one_hot(y, n_classes)

    rate = 0.0001
    pred = SketchNet(x, keep_prob)
    cross_entropy = tf.nn.softmax_cross_entropy_with_logits(logits= pred,
                                                            labels= one_hot_y)

    loss_operation = tf.reduce_mean(cross_entropy)
    optimizer = tf.train.AdamOptimizer(learning_rate=rate)
    training_operation = optimizer.minimize(loss_operation)
    correct_pred = tf.equal(tf.argmax(pred, 1), tf.argmax(one_hot_y, 1))
    accuracy_op = tf.reduce_mean(tf.cast(correct_pred, tf.float32))

    import time
    epoch_accuracy = []
    EPOCHS = 500
    BATCH_SIZE = 64
    save_file = './test_model{0}.ckpt'.format(i)
    saver = tf.train.Saver()
    # print('Saver successfully created')
    with tf.Session() as sess:
        init = tf.global_variables_initializer()
        sess.run(init)
        cur_time = time.time()
        for epoch in range(EPOCHS):
            X_train, y_train = shuffle(X_train, y_train)
            for offset in range(0, n_train, BATCH_SIZE):
                images_feed= X_train[offset:offset+BATCH_SIZE]
                labels_feed = y_train[offset:offset+BATCH_SIZE]
                _, loss_value = sess.run([training_operation, loss_operation],
                                         feed_dict={x: images_feed, y: labels_feed})

            if (epoch+1) % 10 == 0:
                validation_accuracy = evaluate(X_valid, y_valid)

```

```

        epoch_accuracy.append((validation_accuracy, loss_value))
        prev_time = cur_time
        cur_time = time.time()
        print('EPOCH {}'.format(epoch+1))
        print('Loss Value {}'.format(loss_value))
        print('Validation Accuracy = {:.3f}'.format(validation_acc))
        print('Time/epoch {0} s\n'.format((cur_time - prev_time)/1))
        saver.save(sess, save_file)
        print('Saved {0} successfully'.format(save_file))
    epoch_accuracy_loss = []
    epoch_accuracy_acc = []
    for epoch in epoch_accuracy:
        epoch_accuracy_acc.append(epoch[0])
        epoch_accuracy_loss.append(epoch[1])
    epochs = np.arange(10,EPOCHS+1,10)
    accuracy = pd.DataFrame({'Epoch': epochs,
                             'Loss': epoch_accuracy_loss,
                             'Accuracy': epoch_accuracy_acc})

    accuracy
    accuracy.to_csv('accuracy_testmodel{0}.csv'.format(i), index=False)

```

```

Layer 0 Init: Shape (?, 32, 32, 1)
Layer 1 Convolution: Shape (?, 32, 32, 32)
Layer 2 Convolution: Shape (?, 32, 32, 32)
Layer 3 Pooling: Shape (?, 16, 16, 32)
Layer 4 Convolution: Shape (?, 16, 16, 64)
Layer 5 Convolution: Shape (?, 16, 16, 64)
Layer 6 Pooling: Shape (?, 8, 8, 64)
Layer 7 Convolution: Shape (?, 8, 8, 128)
Layer 8 Convolution: Shape (?, 8, 8, 128)
Layer 9 Pooling: Shape (?, 4, 4, 128)
Layer 10 Flatten: Shape (?, 2048)
Layer 11 Matrix Multiplication: Shape (?, 400)
Layer 12 Matrix Multiplication: Shape (?, 100)
Final Layer Matrix Multiplication: Shape (?, 43)

```

```

EPOCH 10
Loss Value 3.490654468536377
Validation Accuracy = 0.084
Time/epoch 53.51769535541534 s

```

```

EPOCH 20
Loss Value 3.058199882507324
Validation Accuracy = 0.156
Time/epoch 53.209389090538025 s

```

```

EPOCH 30
Loss Value 1.5504026412963867
Validation Accuracy = 0.581

```

Time/epoch 53.21339662075043 s

EPOCH 40

Loss Value 0.39222604036331177

Validation Accuracy = 0.816

Time/epoch 53.20665421485901 s

EPOCH 50

Loss Value 0.36341074109077454

Validation Accuracy = 0.870

Time/epoch 53.194975233078004 s

EPOCH 60

Loss Value 0.22096551954746246

Validation Accuracy = 0.872

Time/epoch 53.15903890132904 s

EPOCH 70

Loss Value 0.25180310010910034

Validation Accuracy = 0.882

Time/epoch 53.149676012992856 s

EPOCH 80

Loss Value 0.19409331679344177

Validation Accuracy = 0.901

Time/epoch 53.07203860282898 s

EPOCH 90

Loss Value 0.1360066682100296

Validation Accuracy = 0.910

Time/epoch 53.068830490112305 s

EPOCH 100

Loss Value 0.11865735054016113

Validation Accuracy = 0.913

Time/epoch 53.064456129074095 s

EPOCH 110

Loss Value 0.005096019711345434

Validation Accuracy = 0.901

Time/epoch 53.062315058708194 s

EPOCH 120

Loss Value 0.07833929359912872

Validation Accuracy = 0.917

Time/epoch 53.14083995819092 s

EPOCH 130

Loss Value 0.000202967319637537  
Validation Accuracy = 0.907  
Time/epoch 53.05740141868591 s

EPOCH 140  
Loss Value 0.2923966646194458  
Validation Accuracy = 0.927  
Time/epoch 53.066979217529294 s

EPOCH 150  
Loss Value 0.2569212317466736  
Validation Accuracy = 0.926  
Time/epoch 53.062325620651244 s

EPOCH 160  
Loss Value 0.07807845622301102  
Validation Accuracy = 0.920  
Time/epoch 53.06166276931763 s

EPOCH 170  
Loss Value 0.1436450481414795  
Validation Accuracy = 0.913  
Time/epoch 53.054788041114804 s

EPOCH 180  
Loss Value 0.17633110284805298  
Validation Accuracy = 0.923  
Time/epoch 53.064660000801084 s

EPOCH 190  
Loss Value 0.059870388358831406  
Validation Accuracy = 0.920  
Time/epoch 53.06181490421295 s

EPOCH 200  
Loss Value 0.00027461149147711694  
Validation Accuracy = 0.930  
Time/epoch 53.05368151664734 s

EPOCH 210  
Loss Value 0.23646923899650574  
Validation Accuracy = 0.922  
Time/epoch 53.06076426506043 s

EPOCH 220  
Loss Value 0.12256752699613571  
Validation Accuracy = 0.927  
Time/epoch 53.05485310554504 s



EPOCH 230  
Loss Value 0.11755076050758362  
Validation Accuracy = 0.928  
Time/epoch 53.06219832897186 s

EPOCH 240  
Loss Value 0.11759935319423676  
Validation Accuracy = 0.927  
Time/epoch 53.05841197967529 s

EPOCH 250  
Loss Value 0.11766086518764496  
Validation Accuracy = 0.932  
Time/epoch 53.063020491600035 s

EPOCH 260  
Loss Value 0.059147242456674576  
Validation Accuracy = 0.932  
Time/epoch 53.06206090450287 s

EPOCH 270  
Loss Value 0.11871938407421112  
Validation Accuracy = 0.937  
Time/epoch 53.05417647361755 s

EPOCH 280  
Loss Value 0.11754101514816284  
Validation Accuracy = 0.931  
Time/epoch 53.06199131011963 s

EPOCH 290  
Loss Value 0.09802286326885223  
Validation Accuracy = 0.937  
Time/epoch 53.06087872982025 s

EPOCH 300  
Loss Value 0.17670820653438568  
Validation Accuracy = 0.938  
Time/epoch 53.06172118186951 s

EPOCH 310  
Loss Value 0.11756624281406403  
Validation Accuracy = 0.926  
Time/epoch 53.058766102790834 s

EPOCH 320  
Loss Value 0.23521068692207336

Validation Accuracy = 0.926  
Time/epoch 53.06044363975525 s

EPOCH 330  
Loss Value 0.1778244972229004  
Validation Accuracy = 0.940  
Time/epoch 53.064014768600465 s

EPOCH 340  
Loss Value 0.11777230352163315  
Validation Accuracy = 0.935  
Time/epoch 53.05599873065948 s

EPOCH 350  
Loss Value 0.2350798398256302  
Validation Accuracy = 0.916  
Time/epoch 53.05869874954224 s

EPOCH 360  
Loss Value 0.05903824046254158  
Validation Accuracy = 0.937  
Time/epoch 53.07547943592071 s

EPOCH 370  
Loss Value 0.05878651514649391  
Validation Accuracy = 0.936  
Time/epoch 53.063064742088315 s

EPOCH 380  
Loss Value 1.2647128642129246e-06  
Validation Accuracy = 0.937  
Time/epoch 53.05747218132019 s

EPOCH 390  
Loss Value 0.059153106063604355  
Validation Accuracy = 0.937  
Time/epoch 53.087419271469116 s

EPOCH 400  
Loss Value 0.1763305962085724  
Validation Accuracy = 0.925  
Time/epoch 54.170326948165894 s

EPOCH 410  
Loss Value 0.11753951758146286  
Validation Accuracy = 0.929  
Time/epoch 54.043428134918216 s

```
EPOCH 420
Loss Value 0.41183745861053467
Validation Accuracy = 0.924
Time/epoch 53.63017358779907 s

EPOCH 430
Loss Value 3.4654676710488275e-05
Validation Accuracy = 0.936
Time/epoch 53.38587491512298 s

EPOCH 440
Loss Value 0.1791813224554062
Validation Accuracy = 0.937
Time/epoch 54.32847898006439 s

EPOCH 450
Loss Value 0.23517253994941711
Validation Accuracy = 0.935
Time/epoch 55.025445222854614 s

EPOCH 460
Loss Value 0.12134981155395508
Validation Accuracy = 0.943
Time/epoch 54.64195365905762 s

EPOCH 470
Loss Value 0.1175454631447792
Validation Accuracy = 0.944
Time/epoch 56.21623985767364 s

EPOCH 480
Loss Value 0.004683051258325577
Validation Accuracy = 0.932
Time/epoch 54.38817734718323 s

EPOCH 490
Loss Value 0.05979417636990547
Validation Accuracy = 0.934
Time/epoch 54.567841219902036 s

EPOCH 500
Loss Value 0.05877096578478813
Validation Accuracy = 0.931
Time/epoch 54.53912818431854 s

Saved ./test_model7.ckpt successfully
Layer 0 Init: Shape (?, 32, 32, 1)
Layer 1 Convolution: Shape (?, 32, 32, 32)
```

Layer 2 Convolution: Shape (?, 32, 32, 32)  
Layer 3 Pooling: Shape (?, 16, 16, 32)  
Layer 4 Convolution: Shape (?, 16, 16, 64)  
Layer 5 Convolution: Shape (?, 16, 16, 64)  
Layer 6 Pooling: Shape (?, 8, 8, 64)  
Layer 7 Convolution: Shape (?, 8, 8, 128)  
Layer 8 Convolution: Shape (?, 8, 8, 128)  
Layer 9 Pooling: Shape (?, 4, 4, 128)  
Layer 10 Flatten: Shape (?, 2048)  
Layer 11 Matrix Multiplication: Shape (?, 400)  
Layer 12 Matrix Multiplication: Shape (?, 100)  
Final Layer Matrix Multiplication: Shape (?, 43)  
EPOCH 10  
Loss Value 3.5030837059020996  
Validation Accuracy = 0.062  
Time/epoch 54.39298417568207 s  
  
EPOCH 20  
Loss Value 3.3996071815490723  
Validation Accuracy = 0.069  
Time/epoch 54.2818205833435 s  
  
EPOCH 30  
Loss Value 1.8172221183776855  
Validation Accuracy = 0.437  
Time/epoch 54.44207174777985 s  
  
EPOCH 40  
Loss Value 0.5546555519104004  
Validation Accuracy = 0.685  
Time/epoch 54.010689401626585 s  
  
EPOCH 50  
Loss Value 0.5299626588821411  
Validation Accuracy = 0.768  
Time/epoch 53.23627066612244 s  
  
EPOCH 60  
Loss Value 0.528257429599762  
Validation Accuracy = 0.800  
Time/epoch 53.74606647491455 s  
  
EPOCH 70  
Loss Value 0.381015419960022  
Validation Accuracy = 0.819  
Time/epoch 55.319636654853824 s  
  
EPOCH 80

Loss Value 0.37777870893478394  
Validation Accuracy = 0.830  
Time/epoch 55.099873065948486 s

EPOCH 90  
Loss Value 0.29659515619277954  
Validation Accuracy = 0.835  
Time/epoch 54.47803463935852 s

EPOCH 100  
Loss Value 0.19498413801193237  
Validation Accuracy = 0.836  
Time/epoch 54.02949392795563 s

EPOCH 110  
Loss Value 0.41443532705307007  
Validation Accuracy = 0.842  
Time/epoch 54.18805747032165 s

EPOCH 120  
Loss Value 0.6127446889877319  
Validation Accuracy = 0.867  
Time/epoch 53.72728931903839 s

EPOCH 130  
Loss Value 0.418956995010376  
Validation Accuracy = 0.855  
Time/epoch 53.66988160610199 s

EPOCH 140  
Loss Value 0.0687803104519844  
Validation Accuracy = 0.851  
Time/epoch 53.85901670455932 s

EPOCH 150  
Loss Value 0.24131038784980774  
Validation Accuracy = 0.852  
Time/epoch 55.64516479969025 s

EPOCH 160  
Loss Value 0.3192443251609802  
Validation Accuracy = 0.866  
Time/epoch 53.648840475082395 s

EPOCH 170  
Loss Value 0.37455615401268005  
Validation Accuracy = 0.860  
Time/epoch 54.411151719093326 s

EPOCH 180  
Loss Value 0.1807766854763031  
Validation Accuracy = 0.862  
Time/epoch 53.55290441513061 s

EPOCH 190  
Loss Value 0.2940000891685486  
Validation Accuracy = 0.868  
Time/epoch 54.41455183029175 s

EPOCH 200  
Loss Value 0.11767612397670746  
Validation Accuracy = 0.861  
Time/epoch 53.81856007575989 s

EPOCH 210  
Loss Value 0.2563493251800537  
Validation Accuracy = 0.863  
Time/epoch 54.628816628456114 s

EPOCH 220  
Loss Value 0.18293353915214539  
Validation Accuracy = 0.867  
Time/epoch 56.88869185447693 s

EPOCH 230  
Loss Value 0.3303506374359131  
Validation Accuracy = 0.866  
Time/epoch 54.76349241733551 s

EPOCH 240  
Loss Value 0.06321946531534195  
Validation Accuracy = 0.867  
Time/epoch 55.17261061668396 s

EPOCH 250  
Loss Value 0.17795386910438538  
Validation Accuracy = 0.873  
Time/epoch 56.41953554153442 s

EPOCH 260  
Loss Value 0.2940717041492462  
Validation Accuracy = 0.881  
Time/epoch 55.66902503967285 s

EPOCH 270  
Loss Value 0.6465493440628052

Validation Accuracy = 0.882  
Time/epoch 56.59350936412811 s

EPOCH 280  
Loss Value 0.15138055384159088  
Validation Accuracy = 0.861  
Time/epoch 56.29117460250855 s

EPOCH 290  
Loss Value 0.36115971207618713  
Validation Accuracy = 0.872  
Time/epoch 56.34417695999146 s

EPOCH 300  
Loss Value 0.2579424977302551  
Validation Accuracy = 0.878  
Time/epoch 55.05306560993195 s

EPOCH 310  
Loss Value 0.17779144644737244  
Validation Accuracy = 0.880  
Time/epoch 55.75719492435455 s

EPOCH 320  
Loss Value 0.41795089840888977  
Validation Accuracy = 0.879  
Time/epoch 55.80955755710602 s

EPOCH 330  
Loss Value 0.5322633981704712  
Validation Accuracy = 0.887  
Time/epoch 56.92654228210449 s

EPOCH 340  
Loss Value 0.117564857006073  
Validation Accuracy = 0.878  
Time/epoch 56.66927878856659 s

EPOCH 350  
Loss Value 0.2940155267715454  
Validation Accuracy = 0.885  
Time/epoch 54.7671293258667 s

EPOCH 360  
Loss Value 0.23586395382881165  
Validation Accuracy = 0.879  
Time/epoch 56.00897183418274 s

EPOCH 370  
Loss Value 0.23537227511405945  
Validation Accuracy = 0.883  
Time/epoch 56.41729147434235 s

EPOCH 380  
Loss Value 0.1763276308774948  
Validation Accuracy = 0.883  
Time/epoch 55.15845437049866 s

EPOCH 390  
Loss Value 0.11897905170917511  
Validation Accuracy = 0.883  
Time/epoch 53.79849669933319 s

EPOCH 400  
Loss Value 0.23510581254959106  
Validation Accuracy = 0.880  
Time/epoch 53.63479323387146 s

EPOCH 410  
Loss Value 0.4709963798522949  
Validation Accuracy = 0.872  
Time/epoch 53.64160153865814 s

EPOCH 420  
Loss Value 0.13514569401741028  
Validation Accuracy = 0.881  
Time/epoch 53.6414767742157 s

EPOCH 430  
Loss Value 0.29384535551071167  
Validation Accuracy = 0.888  
Time/epoch 53.710523867607115 s

EPOCH 440  
Loss Value 0.1773061901330948  
Validation Accuracy = 0.888  
Time/epoch 53.762827491760255 s

EPOCH 450  
Loss Value 0.4126024842262268  
Validation Accuracy = 0.893  
Time/epoch 53.76892328262329 s

EPOCH 460  
Loss Value 0.2502591013908386  
Validation Accuracy = 0.889



Time/epoch 53.76353361606598 s

EPOCH 470

Loss Value 0.3526131510734558

Validation Accuracy = 0.875

Time/epoch 53.763596415519714 s

EPOCH 480

Loss Value 0.05878424644470215

Validation Accuracy = 0.889

Time/epoch 53.77308981418609 s

EPOCH 490

Loss Value 0.35280948877334595

Validation Accuracy = 0.886

Time/epoch 53.780203628540036 s

EPOCH 500

Loss Value 0.11754566431045532

Validation Accuracy = 0.884

Time/epoch 53.77831091880798 s

Saved ./test\_model8.ckpt successfully

Layer 0 Init: Shape (?, 32, 32, 1)

Layer 1 Convolution: Shape (?, 32, 32, 32)

Layer 2 Convolution: Shape (?, 32, 32, 32)

Layer 3 Pooling: Shape (?, 16, 16, 32)

Layer 4 Convolution: Shape (?, 16, 16, 64)

Layer 5 Convolution: Shape (?, 16, 16, 64)

Layer 6 Pooling: Shape (?, 8, 8, 64)

Layer 7 Convolution: Shape (?, 8, 8, 128)

Layer 8 Convolution: Shape (?, 8, 8, 128)

Layer 9 Pooling: Shape (?, 4, 4, 128)

Layer 10 Flatten: Shape (?, 2048)

Layer 11 Matrix Multiplication: Shape (?, 400)

Layer 12 Matrix Multiplication: Shape (?, 100)

Final Layer Matrix Multiplication: Shape (?, 43)

EPOCH 10

Loss Value 3.419048309326172

Validation Accuracy = 0.068

Time/epoch 53.90260796546936 s

EPOCH 20

Loss Value 3.2160818576812744

Validation Accuracy = 0.143

Time/epoch 53.934254217147824 s

EPOCH 30

Loss Value 2.082726001739502  
Validation Accuracy = 0.492  
Time/epoch 53.949260401725766 s

EPOCH 40  
Loss Value 1.035753846168518  
Validation Accuracy = 0.744  
Time/epoch 53.93984503746033 s

EPOCH 50  
Loss Value 0.24839964509010315  
Validation Accuracy = 0.840  
Time/epoch 53.9347419500351 s

EPOCH 60  
Loss Value 0.35652652382850647  
Validation Accuracy = 0.873  
Time/epoch 53.929841947555545 s

EPOCH 70  
Loss Value 0.21008503437042236  
Validation Accuracy = 0.889  
Time/epoch 53.938786792755124 s

EPOCH 80  
Loss Value 0.19954434037208557  
Validation Accuracy = 0.900  
Time/epoch 53.919185638427734 s

EPOCH 90  
Loss Value 0.1241820827126503  
Validation Accuracy = 0.914  
Time/epoch 53.92304522991181 s

EPOCH 100  
Loss Value 0.18549934029579163  
Validation Accuracy = 0.909  
Time/epoch 53.95525679588318 s

EPOCH 110  
Loss Value 0.19092458486557007  
Validation Accuracy = 0.920  
Time/epoch 53.95475780963898 s

EPOCH 120  
Loss Value 0.1275901049375534  
Validation Accuracy = 0.924  
Time/epoch 53.96477575302124 s

EPOCH 130  
Loss Value 0.07133034616708755  
Validation Accuracy = 0.924  
Time/epoch 53.9582515001297 s

EPOCH 140  
Loss Value 0.19048258662223816  
Validation Accuracy = 0.917  
Time/epoch 53.96227235794068 s

EPOCH 150  
Loss Value 0.059184327721595764  
Validation Accuracy = 0.920  
Time/epoch 53.965414667129515 s

EPOCH 160  
Loss Value 0.19404970109462738  
Validation Accuracy = 0.924  
Time/epoch 53.97174198627472 s

EPOCH 170  
Loss Value 0.12498500943183899  
Validation Accuracy = 0.931  
Time/epoch 53.965019297599795 s

EPOCH 180  
Loss Value 0.18354877829551697  
Validation Accuracy = 0.927  
Time/epoch 53.97208001613617 s

EPOCH 190  
Loss Value 0.1763223558664322  
Validation Accuracy = 0.935  
Time/epoch 53.98572678565979 s

EPOCH 200  
Loss Value 0.1852569878101349  
Validation Accuracy = 0.932  
Time/epoch 53.989094376564026 s

EPOCH 210  
Loss Value 0.0007359905284829438  
Validation Accuracy = 0.932  
Time/epoch 53.99215724468231 s

EPOCH 220  
Loss Value 0.26579487323760986

Validation Accuracy = 0.933  
Time/epoch 53.99198596477508 s

EPOCH 230  
Loss Value 0.0007353522232733667  
Validation Accuracy = 0.939  
Time/epoch 53.99026136398315 s

EPOCH 240  
Loss Value 0.17665749788284302  
Validation Accuracy = 0.935  
Time/epoch 54.006837940216066 s

EPOCH 250  
Loss Value 0.0001312726817559451  
Validation Accuracy = 0.930  
Time/epoch 53.998423194885255 s

EPOCH 260  
Loss Value 0.0627204030752182  
Validation Accuracy = 0.924  
Time/epoch 53.994708156585695 s

EPOCH 270  
Loss Value 0.17631518840789795  
Validation Accuracy = 0.929  
Time/epoch 54.04091284275055 s

EPOCH 280  
Loss Value 0.17656394839286804  
Validation Accuracy = 0.931  
Time/epoch 54.04328784942627 s

EPOCH 290  
Loss Value 0.12049417197704315  
Validation Accuracy = 0.935  
Time/epoch 54.05194664001465 s

EPOCH 300  
Loss Value 0.17631638050079346  
Validation Accuracy = 0.927  
Time/epoch 54.03776979446411 s

EPOCH 310  
Loss Value 0.17631463706493378  
Validation Accuracy = 0.936  
Time/epoch 54.04577875137329 s

EPOCH 320  
Loss Value 0.11758115887641907  
Validation Accuracy = 0.936  
Time/epoch 54.044568634033205 s

EPOCH 330  
Loss Value 0.1763652265071869  
Validation Accuracy = 0.930  
Time/epoch 54.050613379478456 s

EPOCH 340  
Loss Value 0.35262590646743774  
Validation Accuracy = 0.933  
Time/epoch 54.04581396579742 s

EPOCH 350  
Loss Value 0.11753824353218079  
Validation Accuracy = 0.939  
Time/epoch 54.06279170513153 s

EPOCH 360  
Loss Value 0.05877010524272919  
Validation Accuracy = 0.943  
Time/epoch 54.07713875770569 s

EPOCH 370  
Loss Value 0.05882471054792404  
Validation Accuracy = 0.934  
Time/epoch 54.081227922439574 s

EPOCH 380  
Loss Value 0.21368181705474854  
Validation Accuracy = 0.928  
Time/epoch 54.08398704528808 s

EPOCH 390  
Loss Value 0.17630665004253387  
Validation Accuracy = 0.932  
Time/epoch 54.082628560066226 s

EPOCH 400  
Loss Value 0.17663967609405518  
Validation Accuracy = 0.932  
Time/epoch 54.081160116195676 s

EPOCH 410  
Loss Value 0.1763269454240799  
Validation Accuracy = 0.943

Time/epoch 54.087117075920105 s

EPOCH 420

Loss Value 0.1763089895248413

Validation Accuracy = 0.934

Time/epoch 54.0866712808609 s

EPOCH 430

Loss Value 0.1201428472995758

Validation Accuracy = 0.931

Time/epoch 54.080964183807374 s

EPOCH 440

Loss Value 0.1763075590133667

Validation Accuracy = 0.938

Time/epoch 53.98373551368714 s

EPOCH 450

Loss Value 0.1763285994529724

Validation Accuracy = 0.940

Time/epoch 53.47293472290039 s

EPOCH 460

Loss Value 0.11759638786315918

Validation Accuracy = 0.934

Time/epoch 53.47245950698853 s

EPOCH 470

Loss Value 0.23507526516914368

Validation Accuracy = 0.940

Time/epoch 53.468425846099855 s

EPOCH 480

Loss Value 0.0005968312034383416

Validation Accuracy = 0.927

Time/epoch 53.46960098743439 s

EPOCH 490

Loss Value 0.14574527740478516

Validation Accuracy = 0.940

Time/epoch 53.471102714538574 s

EPOCH 500

Loss Value 0.2350994348526001

Validation Accuracy = 0.939

Time/epoch 53.47247228622437 s

Saved ./test\_model9.ckpt successfully


```
In [28]: epoch_accuracy_loss = []
epoch_accuracy_acc = []
for epoch in epoch_accuracy:
    epoch_accuracy_acc.append(epoch[0])
    epoch_accuracy_loss.append(epoch[1])
epochs = np.arange(10,1001,10)
accuracy = pd.DataFrame({'Epoch': epochs,
                        'Loss': epoch_accuracy_loss,
                        'Accuracy': epoch_accuracy_acc})

accuracy
accuracy.to_csv('accuracy_testmodel6.csv', index=False)
```

### 3.3 Validation Accuracy

- All networks but one achieved the accuracy threshold.
- The inaccurate network (8) was trained exactly the same as other networks (7&9)
- A prior model (4) with an architecture CPCPCPFFF was just as accurate as VGGs CCPC-CPCCPFFF architecture

```
In [133]: # Models 6-9 trained with VGG-like net
fig = plt.figure(figsize=(12,6))
fig.suptitle('Validation Accuracy Model'.format(i))
ax1.set_xlim(left=0, right=1000)
ax1.set_ylim(top=1.0, bottom=0.0)
ax1 = plt.subplot(111)
plt.text(0,.95, '93% Threshold', fontsize=16)
ax1.plot((0,1000), (.93,.93), 'k-')
for i in range(4,10):
    if i != 5:
        accuracy = pd.read_csv('accuracy_testmodel{0}.csv'.format(i))
        ax1.plot(accuracy['Epoch'], accuracy['Accuracy'], label='Model {0}'.format(i))
legend = ax1.legend()
plt.show()
```



### 3.4 Ensemble

Ensemble bagging across 4 networks was carried out. It didn't have any impact. This is probably because usually 10 networks are used in ensemble learning and the networks are only different in their initialization and EPOCH numbers. ## Restoring models \* Major lesson learned. restore expects a path so use ./ if in root directory... \* Test set accuracy is similar to validation set accuracy so overfitting to the validation set is not an issue.

```
In [97]: with tf.Session() as sess:
    BATCH_SIZE = 64
    num_examples = len(X_valid)
    saver = tf.train.Saver()

    sess.run(tf.global_variables_initializer())
    saver.restore(sess, './test_model6.ckpt')
    print('Model 6 Validation Accuracy: {}'.format(evaluate(X_valid, y_val)))
    print('Model 6 Test Accuracy: {}'.format(evaluate(X_test, y_test, sess)))
    pred_valid_6 = predict_labels(X_valid)
    pred_test_6 = predict_labels(X_test)

    sess2 = tf.Session()
    saver.restore(sess2, './test_model7.ckpt')
    print('Model 7 Validation Accuracy: {}'.format(evaluate(X_valid, y_val)))
    print('Model 7 Test Accuracy: {}'.format(evaluate(X_test, y_test, sess2)))
    pred_valid_7 = predict_labels(X_valid)
    pred_test_7 = predict_labels(X_test)
```



```

sess3 = tf.Session()
saver.restore(sess3, './test_model9.ckpt')
print('Model 9 Validation Accuracy: {}'.format(evaluate(X_valid, y_val
print('Model 9 Test Accuracy: {}'.format(evaluate(X_test, y_test, sess
pred_valid_9 = predict_labels(X_valid)
pred_test_9 = predict_labels(X_test)

sess4 = tf.Session()
saver.restore(sess4, './test_model8.ckpt')
print('Model 8 Validation Accuracy: {}'.format(evaluate(X_valid, y_val
print('Model 8 Test Accuracy: {}'.format(evaluate(X_test, y_test, sess
pred_valid_8 = predict_labels(X_valid)
pred_test_8 = predict_labels(X_test)

ensemble_pred = tf.divide(tf.add_n([pred_valid_6, pred_valid_7, pred_v
ensemble_x_entropy = sess.run(tf.nn.softmax_cross_entropy_with_logits

                                feed_dict={y: y_valid})
ensemble_loss = sess.run(tf.reduce_mean(ensemble_x_entropy))
ensemble_correct_pred = sess.run(tf.equal(tf.argmax(ensemble_pred, 1),
                                feed_dict={y: y_valid})
ensemble_accuracy= sess.run(tf.reduce_mean(tf.cast(ensemble_correct_pr
print('Ensemble Accuracy: {0}'.format(ensemble_accuracy))

ensemble_test = tf.divide(tf.add_n([pred_test_6, pred_test_7, pred_te
ensemble_x_entropy = sess.run(tf.nn.softmax_cross_entropy_with_logits

                                feed_dict={y: y_test})
ensemble_loss = sess.run(tf.reduce_mean(ensemble_x_entropy))
ensemble_correct_pred = sess.run(tf.equal(tf.argmax(ensemble_test, 1),
                                feed_dict={y: y_test})
ensemble_accuracy= sess.run(tf.reduce_mean(tf.cast(ensemble_correct_pr
print('Ensemble Accuracy Test: {0}'.format(ensemble_accuracy))
#     loss_operation = tf.reduce_mean(cross_entropy)
#     optimizer = tf.train.AdamOptimizer(learning_rate=rate)
#     training_operation = optimizer.minimize(loss_operation)
#     correct_pred = tf.equal(tf.argmax(pred, 1), tf.argmax(one_hot_y, 1))
#     accuracy_op = tf.reduce_mean(tf.cast(correct_pred, tf.float32))

```

```

Model 6 Validation Accuracy: 0.9498866210719085
Model 6 Test Accuracy: 0.9447347585303578
Model 7 Validation Accuracy: 0.93129251679055
Model 7 Test Accuracy: 0.9383214568959657
Model 9 Validation Accuracy: 0.9394557815560408
Model 9 Test Accuracy: 0.9279493270652783
Model 8 Validation Accuracy: 0.8836734694688498
Model 8 Test Accuracy: 0.8868566904573901
Ensemble Accuracy: 0.9498866200447083

```

Ensemble Accuracy Test: 0.9447347521781921

```
In [105]: def predict_labels(X_data):
            num_examples = len(X_data)
            total_accuracy = 0
            sess = tf.get_default_session()
            all_predictions = sess.run(pred, feed_dict={x: X_data[0:BATCH_SIZE],
            for offset in range(BATCH_SIZE, num_examples, BATCH_SIZE):
                batch_x = X_data[offset:offset+BATCH_SIZE]
                batch_predictions = sess.run(pred, feed_dict={x: batch_x, keep_pr
                all_predictions = tf.concat([all_predictions, batch_predictions],
            return all_predictions
```

### 3.5 Acquiring New Images

- Images were found by driving around with Google streetview in Hamburg.
- The images from streetview were skewed from looking at signs from angle
- The fourth image is close to the ground and is more skewed than the others.
- The background of the fifth image is a truck cabin.
- These images should be easy for the model to predict
- In order to be tested the images needed to be resized to 32x32 and normalized

```
In [103]: from PIL import Image
```

```
In [123]: for i in range(9,15):
            img = cv2.imread('./examples/Selection_{0:03d}.png'.format(i))
            resized = cv2.resize(img, (32,32), 0, 0, cv2.INTER_CUBIC)
            cv2.imwrite('./examples/Resized_{0:03d}.png'.format(i), resized)
            # resized = cv2.cvtColor(resized, cv2.COLOR_BGR2RGB)
            # plt.subplot(230+i)
            # plt.imshow(resized)
            images = np.zeros((6,32,32))
            fig = plt.figure(figsize=(10,9))
            fig.suptitle('Extra images')
            for i in range(6):
                img = mpimg.imread('./examples/Resized_{0:03d}.png'.format(i+9))
                ax = plt.subplot(4,3,i+1)
                plt.imshow(img)
                gray = cv2.cvtColor(img, cv2.COLOR_RGB2GRAY)
                norm = normalize(gray)
                axn = plt.subplot(4,3,i+7)
                plt.imshow(norm, cmap='gray')
                images[i] = norm
            plt.show()
```

## Extra images



```
In [132]: a = []
          with tf.Session() as sess:
              saver = tf.train.Saver()
              sess.run(tf.global_variables_initializer())
              saver.restore(sess, './test_model6.ckpt')
              test_accuracy = evaluate(X_test, y_test, sess)
              print('Test Accuracy is {:.2f}%'.format(test_accuracy*100))
              guess = sess.run(pred, feed_dict={x: images, keep_prob: 1.0})
              smax = tf.nn.softmax(guess)
              top_5 = sess.run(tf.nn.top_k(smax, k=5, sorted=True))
```

```

fig = plt.figure(figsize=(15,15))
fig.suptitle('Extra images')
for i in range(6):
    df5 = pd.DataFrame({'Probability': top_5.values[i],
                        'Sign Name': sign_names['SignName'][top_5.indices[i]]})
    df5 = df5.set_index('Sign Name')
    print(df5, '\n')
    ax1 = plt.subplot(6,2,i*2+2)
    ax2 = plt.subplot(6,2,i*2+1)
    ax1.imshow(images[i], cmap='gray')
    df5.plot.barh(ax=ax2)
plt.show()
#     print(pd.DataFrame({'a': top_5.indices, 'b': top_5.values}))
#     df5 = pd.DataFrame({'SignIndex': top_5.indices,
#                         'Probabilities': top_5.values})
#     for i in range(images.shape[0]):
#         print(guess[i])
#         print(sign_names['SignName'][guess[i].argsort()[0:5]])

```

Test Accuracy is 94.47%

	Probability
Sign Name	
Priority road	1.000000e+00
Roundabout mandatory	1.093382e-10
Speed limit (20km/h)	1.289648e-12
Speed limit (30km/h)	1.289648e-12
Speed limit (50km/h)	1.289648e-12

	Probability
Sign Name	
Go straight or right	0.699499
Stop	0.273761
Priority road	0.014698
Roundabout mandatory	0.002172
No vehicles	0.000299

	Probability
Sign Name	
General caution	9.987790e-01
Traffic signals	1.165736e-03
Stop	3.694654e-05
Keep right	1.032009e-06
Beware of ice/snow	8.811696e-07

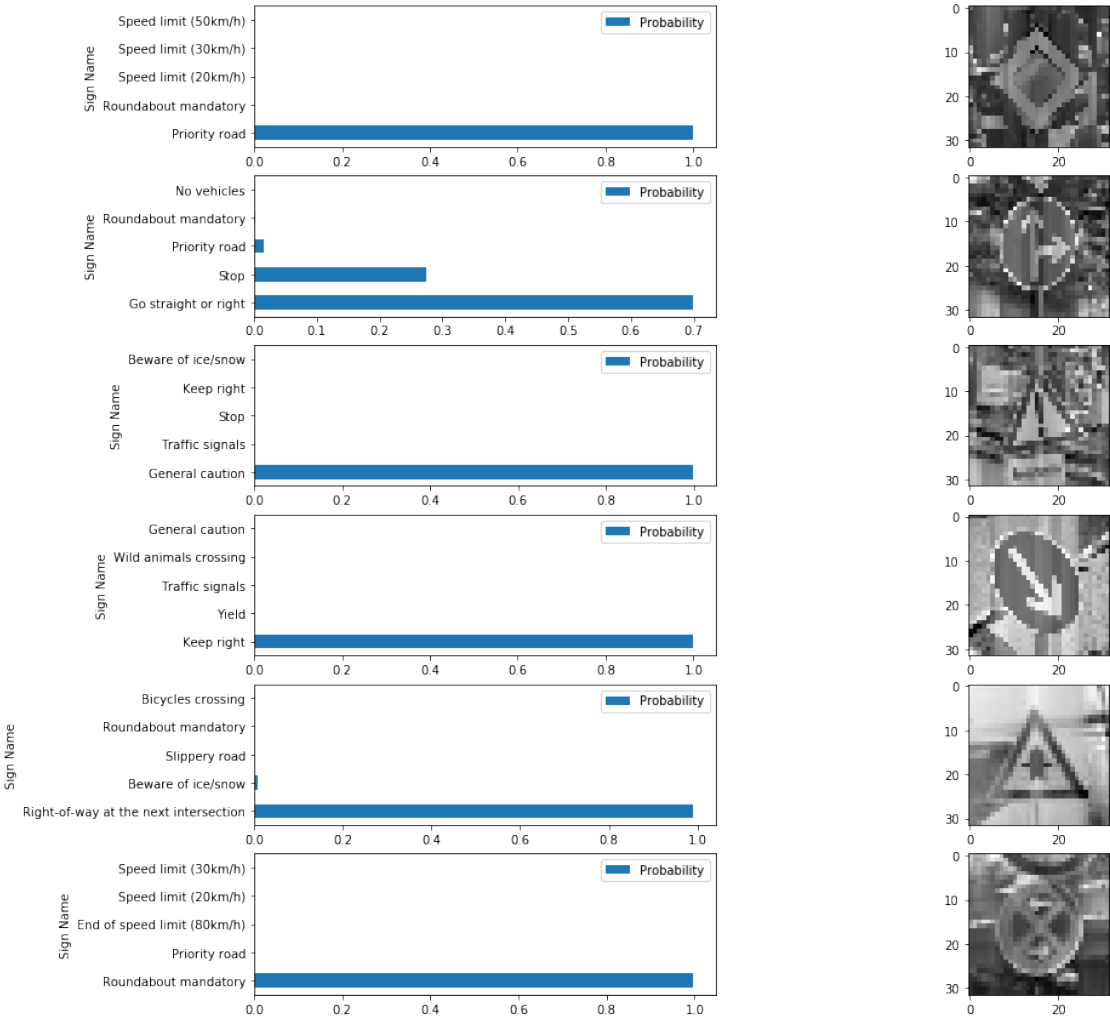
	Probability
Sign Name	
Keep right	1.000000e+00
Yield	1.143891e-26

Traffic signals	1.324083e-29
Wild animals crossing	3.446154e-30
General caution	1.264167e-30

Sign Name	Probability
Right-of-way at the next intersection	9.911767e-01
Beware of ice/snow	8.823336e-03
Slippery road	2.569768e-09
Roundabout mandatory	7.904444e-20
Bicycles crossing	2.305340e-20

Sign Name	Probability
Roundabout mandatory	0.997253
Priority road	0.000694
End of speed limit (80km/h)	0.000371
Speed limit (20km/h)	0.000042
Speed limit (30km/h)	0.000042

Extra images



Successfully classified 6/6 images